


# Logistic Regression





# Logistic Regression


Goal: find the **best line separating** two sets of points





# Logistic Regression

Goal: find the **best line separating** two sets of points





# Optimization Problem

Maximize

$$\sum_{i=1}^n y_i \cdot \log\left(\frac{1}{1 + e^{-(\beta_0 + x_i \cdot \beta)}}\right) + \sum_{i=1}^n (1 - y_i) \cdot \log\left(1 - \frac{1}{1 + e^{-(\beta_0 + x_i \cdot \beta)}}\right)$$

Gradient descent method

$$\beta^{(t+1)} \leftarrow \beta^{(t)} + \alpha \sum_{i=1}^n \left( y_i - \frac{1}{1 + e^{-(\beta_0 + x_i \cdot \beta)}} \right) x_i$$

Sum of values  
calculated with  
each data points



RDD

# Logistic Regression Code


---

- `val data = spark.textFile(...).map(readPoint).cache()`
- `var w = Vector.random(D)`
- `for (i <- 1 to ITERATIONS) {`
- `val gradient = data.map(p =>`
- `(p.y - 1 / (1 + exp(-(w dot p.x)))) * p.x`
- `).reduce(_ + _)`
- `w += gradient`
- `}`
- `println("Final w: " + w)`


# Logistic Regression

- This is also regression but with targets  $Y=(0,1)$ . I.e. it is classification!
- We will fit a regression function on  $P(Y=1|X)$


*linear regression*



$$y_n = AX_n + b$$




*logistic regression*




$$P(Y_n = 1 | X_n) = f(AX_n + b)$$

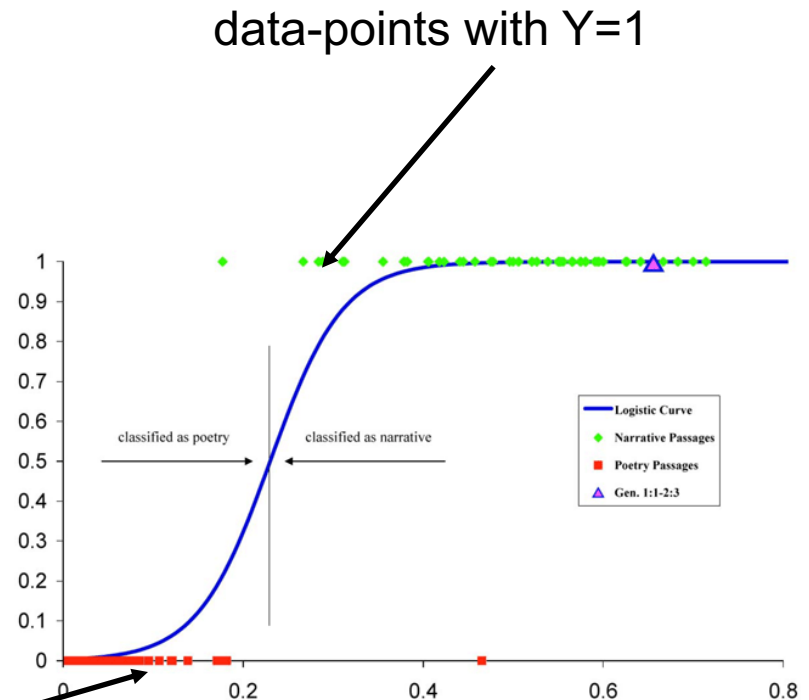
$$f(X) = \frac{1}{1 + \exp[-(AX + b)]}$$



# Sigmoid function $f(x)$




$$P(Y_n = 1 | X_n) = f(AX_n + b)$$



data-points with  $Y=0$

$$f(X) = \frac{1}{1 + \exp[-(AX + b)]}$$

# In 2 Dimensions



A,b determine  
1) orientation  
2) thickness (margin)  
3) offset  
of decision surface

sigmoid  $f(x)$






# Cost Function


We want a different error measure that is better suited for 0/1 data.

This can be derived from maximizing the probability of the data again.

$$P(Y_n = 1 | X_n) = f(AX_n + b)$$

$$P(Y_n = 0 | X_n) = 1 - f(AX_n + b)$$


$$P(Y_n | X_n) = f(X_n)^{Y_n} (1 - f(X_n))^{1-Y_n}$$


$$Error = -\sum_{n=1}^N Y_n \log f(X_n) + (1 - Y_n) \log(1 - f(X_n))$$



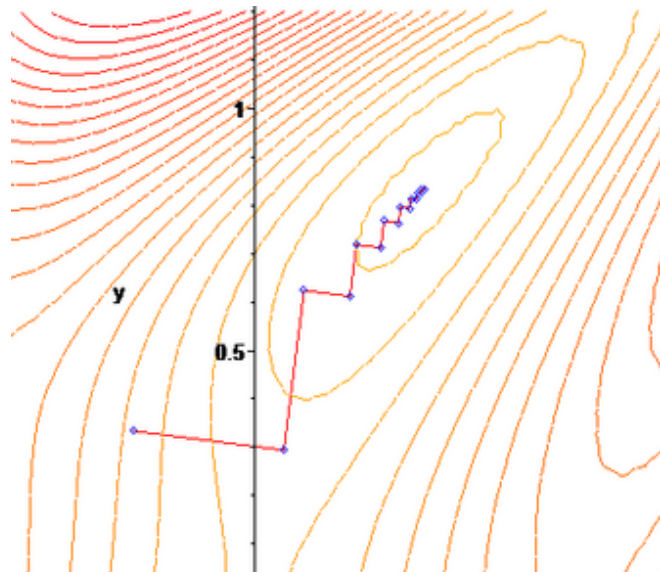

# Learning A,b

Again, we take the derivatives of the Error w.r.t the parameters.

This time however, we can't solve them analytically, so we use gradient descent.

$$A \leftarrow A - \eta \frac{dError}{dA}$$

$$b \leftarrow b - \eta \frac{dError}{db}$$





# Gradients for Logistic Regression

---

After the math (on the white-board) we find:

$$\frac{\partial \text{Error}}{\partial A} = - \sum_n \left[ y_n (1 - f(X_n)) + (1 - y_n) f(X_n) \right] X_n^T$$

$$\frac{\partial \text{Error}}{\partial b} = - \sum_n y_n (1 - f(X_n)) + (1 - y_n) f(X_n)$$

Note: first term in each eqn. (multiplied by Y) only sums over data with Y=1, while second term (multiplied by (1-Y)) only sums over data with Y=0.

Follow the gradient until the change in A,b falls below a small threshold (e.g. 1E-6).



# Classification

---

Once we have found the optimal values for  $A$ ,  $b$  we classify future data with:

$$y_{new} = \text{round}(f(X_{new}))$$

- Least squares and Logistic regression are parametric methods since all the information in the data is stored in the parameters  $A, b$ , i.e. after learning you can toss out the data.
- Also, the decision surface is always linear, its complexity does not grow with the amount of data.
- We have imposed our prior knowledge that the decision surface should be linear.



# A logistic regression learning using TensorFlow

```
import tensorflow as tf
import numpy as np
from sklearn import datasets

iris = datasets.load_iris()
iris_X = iris.data
iris_y = iris.target

np.random.seed(0)
indices = np.random.permutation(len(iris_X))
train_X = iris_X[indices[:-10]]
train_Y = np.asarray([ [ 1 if idx == 0 else 0 for idx in iris_y[indices[:-10]] ] ]).T
test_X = iris_X[indices[-10:]]
test_Y = np.asarray([ [ 1 if idx == 0 else 0 for idx in iris_y[indices[-10:]] ] ]).T
```



# A logistic regression learning using TensorFlow

```
# data format is as usual:
# train_X and test_X have shape (num_instances, num_features)
# train_Y and test_Y have shape (num_instances, num_classes)
num_features = train_X.shape[1]
num_classes = train_Y.shape[1]

# Create variables
# X is a symbolic variable which will contain input data
# shape [None, num_features] suggests that we don't limit the number of
instances in the model
# while the number of features is known in advance
X = tf.placeholder("float", [None, num_features])
# same with labels: number of classes is known, while number of instances is
left undefined
Y = tf.placeholder("float", [None, num_classes])

# W - weights array
W = tf.Variable(tf.zeros([num_features, num_classes]))
# B - bias array
B = tf.Variable(tf.zeros([num_classes]))
```



# A logistic regression learning using TensorFlow

```
# Define a model
# a simple linear model  $y=wx+b$  wrapped into softmax
pY = tf.nn.softmax(tf.matmul(X, W) + B)
# pY will contain predictions the model makes, while Y contains real data

# Define a cost function
cost_fn = -tf.reduce_sum(Y * tf.log(pY))

# Define an optimizer
opt = tf.train.AdamOptimizer(0.01).minimize(cost_fn)
```



# A logistic regression learning using TensorFlow

```
# Create and initialize a session
sess = tf.Session()
init = tf.global_variables_initializer()
sess.run(init)

# run an optimization step with all train data
sess.run(opt, feed_dict={X:train_X, Y:train_Y})

# Now assess the model
# create a variable which reflects how good your predictions are
# here we just compare if the predicted label and the real label are the
same
accuracy = tf.reduce_mean(tf.cast(tf.equal(tf.argmax(pY,1), tf.argmax(Y,1)),
"float"))
# and finally, run calculations with all test data
print(sess.run(accuracy, feed_dict={X:test_X, Y:test_Y}))
```





# A logistic regression learning using Scikit-learn

```
from sklearn import linear_model, datasets

iris = datasets.load_iris()
iris_X = iris.data
iris_y = iris.target

np.random.seed(0)
indices = np.random.permutation(len(iris_X))
train_X = iris_X[indices[:-10]]
train_Y = [ 1 if idx == 0 else 0 for idx in iris_y[indices[:-10]] ]
test_X = iris_X[indices[-10:]]
test_Y = [ 1 if idx == 0 else 0 for idx in iris_y[indices[-10:]] ]

logistic = linear_model.LogisticRegression(C=1e5)
logistic.fit(train_X, train_Y)
print(np.mean([ 1 if a == b else 0 for a, b in zip(logistic.predict(test_X),
test_Y) ]))
```